

# Project 4

Kevin Siraki

April 20, 2021

## 1 Introduction

For this project, we had a similar problem to tackle as in Project 3: making functions run in parallel and measuring the overall run-time of both the original serial code and the parallel code. We began by running the original code with "gettimeofday()" in our main method in order to see the serial run-time. We then proceeded to convert the code to parallel in 3 different ways depending on the type of loop/iterator data type we were dealing with in each case. The results of this experimentation process will be further discussed with in depth code analysis in the following subsections.

## 2 Numerical Integration Serial:

	Serial:	Parallel:	
Small:	0.001001	0.004004	
Medium:	0.250227	0.041038	
Large:	9.115155	2.22297	
			X: Approach Avg. Run-time (Seconds)
			Y: Parameter Size

Numerical Integration featured 5 different techniques of integration. Each of these techniques were split into functions that took in the boundaries of integration, precision, and a function pointer as parameters. The important parameter we manipulated to test run-times was the precision parameter. As the above graph indicates, the serial method had an easy time with small-to-medium precision when calculating integrals, but suffers when we aim for a more precise calculation. These parameter sizes can be seen within the comments on lines 170-177. The corresponding run-times are displayed according to parameter size in the above graphical representation.

---

```
1 #include <stdio.h>
2 #include <stdlib.h>
3 #include <math.h>
4 #include <unistd.h>
5 #include <float.h>
6 #include <sys/time.h>
7
8 double int_leftrect(double from, double to, double n, double (*func)())
9 {
10     double h = (to-from)/n;
11     double sum = 0.0, x;
12     for(x=from; x <= (to-h); x += h)
13         sum += func(x);
14     return h*sum;
15 }
```

```

16
17 double int_rightrect(double from, double to, double n, double (*func)())
18 {
19     double h = (to-from)/n;
20     double sum = 0.0, x;
21     for(x=from; x <= (to-h); x += h)
22         sum += func(x+h);
23     return h*sum;
24 }
25
26 double int_midrect(double from, double to, double n, double (*func)())
27 {
28     double h = (to-from)/n;
29     double sum = 0.0, x;
30     for(x=from; x <= (to-h); x += h)
31         sum += func(x+h/2.0);
32     return h*sum;
33 }
34
35 double int_trapezium(double from, double to, double n, double (*func)())
36 {
37     double h = (to - from) / n;
38     double sum = func(from) + func(to);
39     int i;
40     for(i = 1; i < n; i++)
41         sum += 2.0*func(from + i * h);
42     return h * sum / 2.0;
43 }
44
45 double int_simpson(double from, double to, double n, double (*func)())
46 {
47     double h = (to - from) / n;
48     double sum1 = 0.0;
49     double sum2 = 0.0;
50     int i;
51
52     double x;
53
54     for(i = 0; i < n; i++)
55         sum1 += func(from + h * i + h / 2.0);
56
57     for(i = 1; i < n; i++)
58         sum2 += func(from + h * i);
59
60     return h / 6.0 * (func(from) + func(to) + 4.0 * sum1 + 2.0 * sum2);
61 }
62 /* test */
63 double f3(double x)
64 {
65     return x;
66 }
67
68 double f3a(double x)
69 {

```

```

70     return x*x/2.0;
71 }
72
73 double f2(double x)
74 {
75     return x*x;
76 }
77
78 double f2a(double x)
79 {
80     return x*x*x/3.0;
81 }
82
83 double f1(double x)
84 {
85     return x*x*x;
86 }
87
88 double f1a(double x)
89 {
90     return x*x*x*x/4.0;
91 }
92
93 typedef double (*pfunc)(double, double, double, double (*)());
94 typedef double (*rfunc)(double);
95
96 #define INTG(F,A,B) (F((B))-F((A)))
97
98 int main()
99 {
100     struct timeval start, end;
101
102     gettimeofday(&start, NULL);
103     int i, j;
104     double ic;
105
106     pfunc f[5] = {
107         int_leftrect, int_rightrect,
108         int_midrect, int_trapezium,
109         int_simpson
110     };
111     const char *names[5] = {
112         "leftrect", "rightrect", "midrect",
113         "trapezium", "simpson"
114     };
115     rfunc rf[] = { f1, f2, f3, f3 };
116     rfunc lf[] = { f1a, f2a, f3a, f3a };
117     double ival[] = {
118         0.0, 1.0,
119         1.0, 100.0,
120         0.0, 5000.0,
121         0.0, 6000.0
122     };
123     // double approx[] = { 100.0, 100.0, 100.0, 100.0 }; //very small

```

```

124         //double approx[] = { 1000.0, 1000.0, 5000000.0, 6000000.0 }; //medium
125     double approx[] = { //large
126         99999999,
127         99999999,
128         99999999,
129         99999999
130     };
131     for(j=0; j < (sizeof(rf) / sizeof(rfunc)); j++)
132     {
133         for(i=0; i < 5 ; i++)
134         {
135             ic = (*f[i])(ivals[2*j], ival[2*j+1], approx[j], rf[j]);
136             printf("%10s [ 0,1] num: %+1f, an: %1f\n",
137                 names[i], ic, INTG((*If[j]), ival[2*j], ival[2*j+1]));
138         }
139         printf("\n");
140     }
141     gettimeofday(&end, NULL);
142     double time_taken = end.tv_sec + end.tv_usec / 1e6 -
143         start.tv_sec - start.tv_usec / 1e6; // in seconds
144
145
146     printf("Integration took %f seconds to execute in Serial\n", time_taken);
147 }

```

---

### 3 Numerical Integration Parallel:

	Serial:	Parallel:	
Small:	0.001001	0.004004	
Medium:	0.250227	0.041038	
Large:	9.115155	2.22297	
			X: Approach Avg. Run-time (Seconds)
			Y: Parameter Size

The additions made to the original code to allow this program to run properly in parallel were quite meticulously implemented. To begin with, we utilized fork-and-join with a parallel section in the first 3 integration methods due to the fact that a simple reduction would not suffice. This is due to the fact that the iterator in these functions is a double and the OMP directives within the compiler cannot deal with floating point numbers. Moreover, a simple cast of the iterator to an integer would not be possible because we would be losing enormous amounts of accuracy in our calculations. With this approach, we were able to "fork" the work onto multiple threads, allowing the loops themselves to run in "serial", then be "joined" in a critical section where their total sums were aggregated. The final two integration functions were quite simple to convert to parallel as they only featured simple for loops that had an integer iterator. We converted these to parallel by using a simple work-share construct and reduction. The benefits of the parallel approach were quite apparent as the problem size was increased, as can be seen in the graph above. For a large problem size, the runtime was about 1/4 that of the serial approach. The small and medium problem sizes were still within a margin of error to the serial approach, so this program seems to benefit greatly from being ran in parallel when a very large parameter size is introduced.

```

1 #include <stdio.h>
2 #include <stdlib.h>
3 #include <math.h>
4 #include <omp.h>

```

```

5  #include <float.h>
6  #include <unistd.h>
7  #include <sys/time.h>
8
9  double int_leftrect(double from, double to, double n, double( * func)()) {
10     double h = (to - from) / n;
11     double sum = 0.0, x;
12     double tstart;
13     int tid;
14     int nthreads;
15     double tend;
16     double tsum = 0.0;
17     #pragma omp parallel shared(from, to, n, nthreads, sum) private(x, tid, tsum, tstart, tend)
18     {
19         tid = omp_get_thread_num();
20         nthreads = omp_get_num_threads();
21         tstart = tid * (int) ceil(((double)(to - h) / nthreads));
22         tend = (tid + 1) * (int) ceil(((double)(to - h) / nthreads));
23         if (tend <= (to - h))
24             for (x = tstart; x <= tend; x += h)
25                 tsum += func(x);
26         else
27             for (x = tstart; x <= (to - h); x += h)
28                 tsum += func(x);
29         #pragma omp critical
30             sum += tsum;
31     }
32     return h * sum;
33 }
34
35 double int_rightrect(double from, double to, double n, double( * func)()) {
36     double h = (to - from) / n;
37     double sum = 0.0, x;
38     double tstart;
39     int tid;
40     double tend;
41     int nthreads;
42     double tsum = 0.0;
43     #pragma omp parallel shared(from, to, n, nthreads, sum) private(x, tid, tsum, tstart, tend)
44     {
45         tid = omp_get_thread_num();
46         nthreads = omp_get_num_threads();
47         tstart = tid * (int) ceil(((double)(to - h) / nthreads));
48         tend = (tid + 1) * (int) ceil(((double)(to - h) / nthreads));
49         if (tend < (to - h))
50             for (x = tstart; x <= tend; x += h)
51                 tsum += func(x + h);
52         else
53             for (x = tstart; x <= to - h; x += h)
54                 tsum += func(x + h);
55         #pragma omp critical
56             sum += tsum;
57     }
58     return (h * sum);

```

```

59 }
60
61 double int_midrect(double from, double to, double n, double( * func)()) {
62     double h = (to - from) / n;
63     double sum = 0.0, x;
64     double tstart;
65     int tid;
66     double tend;
67     int nthreads;
68     double tsum = 0.0;
69     #pragma omp parallel shared(from, to, n, sum, nthreads) private(x, tid, tsum, tstart, tend)
70     {
71         tid = omp_get_thread_num();
72         nthreads = omp_get_num_threads();
73         tstart = tid * (int) ceil(((double)(to - h) / nthreads));
74         tend = (tid + 1) * (int) ceil(((double)(to - h) / nthreads));
75         if (tend < (to - h))
76             for (x = tstart; x <= tend; x += h)
77                 tsum += func(x + h / 2.0);
78         else
79             for (x = tstart; x <= to - h; x += h)
80                 tsum += func(x + h / 2.0);
81         #pragma omp critical
82         sum += tsum;
83     }
84     return h * sum;
85 }
86
87 double int_trapezium(double from, double to, double n, double( * func)()) {
88     double h = (to - from) / n;
89     double sum = func(from) + func(to);
90     int i;
91     #pragma omp parallel
92     {
93         #pragma omp for reduction(+: sum)
94         for (i = 1; i < (int) n; i++) {
95             sum += 2.0 * func(from + i * h);
96         }
97     }
98     return h * sum / 2;
99 }
100
101 double int_simpson(double from, double to, double n, double( * func)()) {
102     double h = (to - from) / n;
103     double sum1 = 0.0;
104     double sum2 = 0.0;
105     int i;
106     #pragma omp parallel
107     {
108         #pragma omp for reduction(+: sum1)
109         for (i = 0; i < (int) n; i++) {
110             sum1 += func(from + h * i + h / 2);
111         }
112         #pragma omp for reduction(+: sum2)

```

```

113         for (i = 1; i < (int) n; i++) {
114             sum2 += func(from + h * i);
115         }
116     }
117     return h / 6 * (func(from) + func(to) + 4 * sum1 + 2 * sum2);
118 }
119
120 /* test */
121 double f3(double x) {
122     return x;
123 }
124
125 double f3a(double x) {
126     return x * x / 2.0;
127 }
128
129 double f2(double x) {
130     return x * x ;
131 }
132
133 double f2a(double x) {
134     return x * x * x / 3.0;
135 }
136
137 double f1(double x) {
138     return x * x * x;
139 }
140
141 double f1a(double x) {
142     return x * x * x * x / 4.0;
143 }
144
145 typedef double( * pfunc)(double, double, double, double( * )());
146 typedef double( * rfunc)(double);
147
148 #define INTG(F, A, B)(F((B)) - F((A)))
149
150 int main() {
151     struct timeval start, end;
152
153     gettimeofday( & start, NULL);
154     int i, j;
155     double ic;
156
157     pfunc f[5] = {
158         int_leftrect,
159         int_rightrect,
160         int_midrect,
161         int_trapezium,
162         int_simpson
163     };
164     const char * names[5] = {
165         "leftrect",
166         "rightrect",

```

```

167     "midrect",
168     "trapezium",
169     "simpson"
170 };
171 rfunc rf[] = {
172     f1,
173     f2,
174     f3,
175     f3
176 };
177 rfunc If[] = {
178     f1a,
179     f2a,
180     f3a,
181     f3a
182 };
183 double ivals[] = {
184     0.0,
185     1.0,
186     1.0,
187     100.0,
188     0.0,
189     5000.0,
190     0.0,
191     6000.0
192 };
193
194     // double approx[] = { 100.0, 100.0, 100.0, 100.0 }; //very small
195     //double approx[] = { 1000.0, 1000.0, 5000000.0, 6000000.0 }; //medium
196 double approx[] = { //large
197     99999999,
198     99999999,
199     99999999,
200     99999999
201 };
202 for (j = 0; j < (sizeof(rf) / sizeof(rfunc)); j++) {
203     for (i = 0; i < 5; i++) {
204         ic = ( * f[i])(ivals[2 * j], ivals[2 * j + 1], approx[j], rf[j]);
205         printf("%10s [ 0,1] num: %+1f, an: %1f\n",
206             names[i], ic, INTG(( * If[j]), ivals[2 * j], ivals[2 * j + 1]));
207     }
208     printf("\n");
209 }
210 gettimeofday( & end, NULL);
211 double time_taken = end.tv_sec + end.tv_usec / 1e6 -
212     start.tv_sec - start.tv_usec / 1e6; // in seconds
213
214 printf("Integration took %f seconds to execute in Parallel\n", time_taken);
215 }

```

---



## 4 Runge-Kutta Serial:

	Serial:	Parallel:	
Small:	0.001001	0.001001	
Medium:	0.005003	0.005005	
Large:	4.93593	4.14735	
			X: Approach Avg. Run-time (Seconds)
			Y: Parameter Size

According to Wikipedia, Runge-Kutta is a "In numerical analysis, the Runge–Kutta methods are a family of implicit and explicit iterative methods, which include the well-known routine called the Euler Method, used in temporal discretization for the approximate solutions of ordinary differential equations. " As we can see, the serial method for this function seems to run quite quickly for anything from small to large problem sizes. The size of the problems can be seen in lines 43-47. The parameter size on line 43 was quite small and resulted in too tiny of a run-time to be relevant in our observations. The corresponding run-times based on parameter sizes is present in the chart on top.

---

```

1 // C program to implement Runge
2 // Kutta method
3 #include <stdio.h>
4 #include <unistd.h>
5 #include <sys/time.h>
6 // A sample differential equation
7 // "dy/dx = (x - y)/2"
8
9 float dydx(float x, float y)
10 {
11     return (x + y - 2);
12 }
13
14 // Finds value of y for a given x
15 // using step size h
16 // and initial value y0 at x0.
17 float rungeKutta(float x0, float y0, float x, float h)
18 {
19     // Count number of iterations
20     // using step size or
21     // step height h
22     int n = (int)((x - x0) / h), i;
23     float k1, k2;
24     // Iterate for number of iterations
25     float y = y0;
26     for (i = 1; i <= n; i++)
27     {
28         k1 = h * dydx(x0, y);
29         k2 = h * dydx(x0 + 0.5 * h, y + 0.5 * k1);
30         // Update next value of y
31         y = y + (1.0 / 6.0) * (k1 + 2 * k2);
32         // Update next value of x
33         x0 = x0 + h;
34     }
35     return y;
36 }
37

```

```

38 // Driver Code
39 int main()
40 {
41     struct timeval start, end;
42     gettimeofday(&start, NULL);
43     //float x0 = 0, y = 1, x = 2, h = 0.002; //too small
44
45     //float x0 = 0, y = 10, x = 20, h = 0.001; //small
46     //float x0 = 0, y = 10, x = 20, h = 0.0001; //medium
47     float x0 = 0, y = 10, x = 20, h = 0.0000001; //large
48     printf("y(x) = %f", rungeKutta(x0, y, x, h));
49     gettimeofday(&end, NULL);
50     double time_taken = end.tv_sec + end.tv_usec / 1e6 - start.tv_sec - start.tv_usec / 1e6; // in seconds
51     printf("\nRunge-Kutta took %f seconds to execute in Serial\n", time_taken);
52     return 0;
53 }

```

---

## 5 Runge-Kutta Parallel:

	Serial:	Parallel:	
Small:	0.001001	0.001001	
Medium:	0.005003	0.005005	
Large:	4.93593	4.14735	
			X: Approach Avg. Run-time (Seconds)
			Y: Parameter Size

In our run-time analysis of the parallel version of Runge-Kutta, the performance was equivalent to or slightly worse with a smaller parameter size in comparison to the serial version. However, as expected, a larger parameter size for the variable 'h' showed us the true benefits of running this code in parallel. The parallel version of this code was devised by utilizing tasks due to the nature of the for loop in the Runge-Kutta method. A simple reduction would not work here because parallelization of the for loop would result in an incorrect answer when a messy work-share construct was implemented.

---

```

1 #include <stdio.h>
2 #include <unistd.h>
3 #include <sys/time.h>
4 #include <omp.h>
5
6 float dydx(float x, float y)
7 {
8     return (x + y - 2);
9 }
10
11 float rungeKutta(float x0, float y0, float x, float h)
12 {
13     int n = (int)((x - x0) / h), i;
14     float k1, k2;
15     float y = y0;
16     #pragma omp parallel reduction(+:x0)
17     {
18         #pragma omp master
19         {
20             for (i = 1; i <= n; i++)

```

```

21     {
22         #pragma omp task shared(x0,h,y) private(i,k1,k2)
23         {
24             k1 = h * dydx(x0, y);
25             k2 = h * dydx(x0 + 0.5 * h, y + 0.5 * k1);
26             x0 = x0 + h; //x0 += h;
27             y = y + (1.0 / 6.0) * (k1 + 2 * k2);
28         }
29     }
30 }
31 #pragma omp taskwait
32 }
33 return y;
34 }
35
36 int main()
37 {
38     struct timeval start, end;
39     gettimeofday( & start, NULL);
40     //float x0 = 0, y = 1, x = 2, h = 0.002; //too small
41
42     //float x0 = 0, y = 10, x = 20, h = 0.001; //small
43     //float x0 = 0, y = 10, x = 20, h = 0.0001; //medium
44     float x0 = 0, y = 10, x = 20, h = 0.0000001; //large
45     printf("y(x) = %f", rungeKutta(x0, y, x, h));
46     gettimeofday( & end, NULL);
47     double time_taken = end.tv_sec + end.tv_usec / 1e6 - start.tv_sec - start.tv_usec / 1e6; // in seconds
48     printf("\nRunge-Kutta took %f seconds to execute in Parallel\n", time_taken);
49     return 0;
50 }

```

---

## 6 Conclusion:

To conclude, in conducting our analysis on these Numerical Integration and Runge-Kutta in parallel, we have come to the realization that both of these programs seem to benefit quite greatly from being run in parallel when a larger parameter size is introduced. Moreover, even on a small problem size, the parallel versions of these programs ran at nearly the same speed or within the margin of error of the serial versions. Overall, we utilized fork/join, work-share constructs/reductions, and tasks in Open-MP to run these programs in parallel and the behavior of the parallel code seemed superior to the serial, especially in the case of larger problems being solved.